

Linear Regression

Kyungjae Lee, Yunho Choi, Timothy Ha
RLLAB

[Exercise 1] Univariate Linear Regression


- Univariate linear function: $y = w_1x + w_0$, where $w_0, w_1 \in \mathbb{R}$.
- Let $\mathbf{w} = [w_0, w_1]^T$ and define

$$h_{\mathbf{w}}(x) = w_1x + w_0.$$

- **Linear regression:** Given a training set $\{(x_i, y_i) : 1 \leq i \leq N\}$, find $h_{\mathbf{w}}$ that best fits the training set.

$$Loss(h_{\mathbf{w}}) = \sum_{j=1}^N L_2(y_j, h_{\mathbf{w}}(x_j)) = \sum_{j=1}^N (y_j - h_{\mathbf{w}}(x_j))^2 = \sum_{j=1}^N (y_j - (w_1x_j + w_0))^2.$$


- 1-D linear regression problem
 - Given one dimensional data, find the most fitted line
- Goal
 - Construct linear model in tensorflow graph
 - Define loss function
 - Optimize it using gradient descent method



[Exercise 1] Univariate Linear Regression

1D data generation

```
def f1(x,a,b):  
    return a*x+b  
  
a = 5.  
b = 3.  
  
x = np.linspace(-1,1,100)  
y = f1(x,a,b)  
  
n = 10  
x_data = np.linspace(-1,1,n)  
y_data = f1(x_data,a,b) + np.random.randn(n)  
  
plt.plot(x,y,"k-",label="Ground Truth")  
plt.plot(x_data,y_data,"r.",label="Data")  
plt.legend()  
plt.xlim([-1,1])  
plt.ylim([-2,8])  
plt.show()
```



- Step 1
 - Generate 1d noise data from the underlying linear model
 - Plot it

[Exercise 1] Univariate Linear Regression

- Step 2
 - Define placeholder : x_ph, y_ph
 - `tf.placeholder(dtype, shape=(None,))`

x_ph

y_ph

Tensorflow Graph Construction

```
tf.reset_default_graph()

x_ph =
y_ph =

a_hat =
b_hat =

y_pred =

loss =

# define optimizer
optimizer =
train =

print("Tensorflow graph is constructed!!!")
print("x_ph is:")
print("type: "+str(type(x_ph)))
print("shape: "+str(x_ph.shape))
print("a_hat is:")
print("type: "+str(type(a_hat)))
print("shape: "+str(a_hat.shape))
print("y_pred is:")
print("type: "+str(type(y_pred)))
print("shape: "+str(y_pred.shape))
```

[Exercise 1] Univariate Linear Regression

- Step 3
 - Define tensorflow variable : a_hat, b_hat
 - `tf.get_variable(name, shape=(), initializer=?)`
 - `tf.random_normal_initializer()`

a_hat

x_ph

b_hat

y_ph

Tensorflow Graph Construction

```
tf.reset_default_graph()

x_ph =
y_ph =

a_hat =
b_hat =

y_pred =


loss =

# define optimizer
optimizer =
train =

print("Tensorflow graph is constructed!!!")
print("x_ph is:")
print("type: "+str(type(x_ph)))
print("shape: "+str(x_ph.shape))
print("a_hat is:")
print("type: "+str(type(a_hat)))
print("shape: "+str(a_hat.shape))
print("y_pred is:")
print("type: "+str(type(y_pred)))
print("shape: "+str(y_pred.shape))
```

[Exercise 1] Univariate Linear Regression

- Step 4
 - Define linear model : y_{pred}
 - `tf.multiply(x_ph, a_hat) + b_hat`



Tensorflow Graph Construction

```
tf.reset_default_graph()

x_ph =
y_ph =

a_hat =
b_hat =

y_pred =

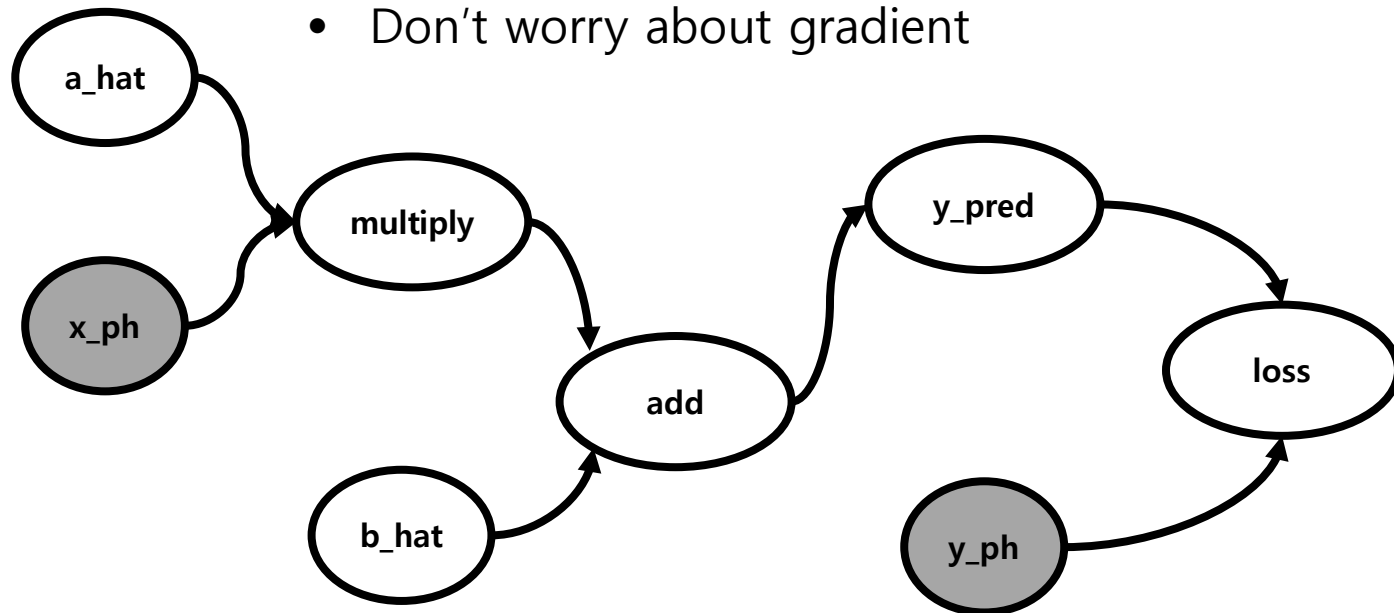
loss =

# define optimizer
optimizer =
train =

print("Tensorflow graph is constructed!!!")
print("x_ph is:")
print("type: "+str(type(x_ph)))
print("shape: "+str(x_ph.shape))
print("a_hat is:")
print("type: "+str(type(a_hat)))
print("shape: "+str(a_hat.shape))
print("y_pred is:")
print("type: "+str(type(y_pred)))
print("shape: "+str(y_pred.shape))
```

[Exercise 1] Univariate Linear Regression

- Step 5
 - Define loss function : loss
 - $0.5 * \text{tf.reduce_mean}(\text{tf.square}(y_{\text{pred}} - y_{\text{ph}}))$
 - Define optimizer : optimizer, train
 - `tf.train.GradientDescentOptimizer(learning_rate=0.1)`
 - `optimizer.minimize(loss)`
 - Don't worry about gradient



Tensorflow Graph Construction

```
tf.reset_default_graph()

x_ph =
y_ph =

a_hat =
b_hat =

y_pred =

loss =

# define optimizer
optimizer =
train =

print("Tensorflow graph is constructed!!!")
print("x_ph is:")
print("type: "+str(type(x_ph)))
print("shape: "+str(x_ph.shape))
print("a_hat is:")
print("type: "+str(type(a_hat)))
print("shape: "+str(a_hat.shape))
print("y_pred is:")
print("type: "+str(type(y_pred)))
print("shape: "+str(y_pred.shape))
```

[Exercise 1] Univariate Linear Regression

- Step 6

- Initialize variable
 - `tf.global_variables_initializer()`
- Open tensorflow session
 - `tf.Session`

Initialize TF variables

```
# initialize variables
init =

sess =
sess.run(init)

a_np, b_np = sess.run([a_hat, b_hat])
print("Randomly initialized parameters")
print("a: %f"%a_np)
print("b: %f"%b_np)
```

- Step 7

- Optimize it!
- Flow the data into placeholder
- Run optimization
- `sess.run([loss, train], feed_dict=feed_dict)`


Optimize linear model (Find the best parameters)

```
nepoch = 3000
feed_dict = {x_ph: x, y_ph: y}
for epoch in range(nepoch):
    loss_np, _ =
    if (epoch%300) == 0:
        print("%d/%d loss : %f"%(epoch, nepoch, loss_np))
    print()

a_np, b_np = sess.run([a_hat, b_hat])
print("Optimized parameters")
print("a: %f"%a_np)
print("b: %f"%b_np)
print("Original parameters")
print("a: %f"%a)
print("b: %f"%b)

feed_dict = {x_ph: x}
y_pred_np = sess.run(y_pred, feed_dict=feed_dict)
plt.plot(x, y, "k-", label="Ground Truth")
plt.plot(x_data, y_data, "r.", label="Data")
plt.plot(x, y_pred_np, "b-", label="Linear Regression")
plt.legend()
plt.xlim([-1, 1])
plt.ylim([-2, 8])
plt.show()
```


[Exercise 2] Multivariate Linear Regression



- Do the same thing for two dimensional data
- Goal
 - Construct linear model in tensorflow graph
 - Define loss function
 - Optimize it using gradient descent method

2D data generation

```
def f2(x,a,b):  
    return np.matmul(x,a)+b  
  
a = np.array([[2.],[1.]])  
b = 1.  
  
x1 = np.linspace(-1,1,100)  
x2 = np.linspace(-1,1,100)  
X1,X2 = np.meshgrid(x1,x2)  
x = np.concatenate([np.reshape(X1,[-1,1]),np.reshape(X2,[-1,1])],axis=1)  
y = f2(x,a,b)  
Y = np.reshape(y,[100,100])  
  
n = 50  
x_data = np.random.uniform(-1,1,size=(n,2))  
y_data = f2(x_data,a,b) + np.random.randn(n,1)  
  
fig = plt.figure()  
ax = fig.gca(projection='3d')  
  
ax.plot_surface(X1,X2,Y)  
ax.plot3D(x_data[:,0].flatten(),x_data[:,1].flatten(),y_data.flatten(),'r.',label="Data")  
plt.legend()  
ax.set_xlim3d(-1,1)  
ax.set_ylim3d(-1,1)  
ax.set_zlim3d(-3,4)  
plt.show()
```

[Exercise 2] Multivariate Linear Regression

- Step 1
 - Define placeholder
 - `tf.placeholder(dtype, shape=(None,2))`
 - `tf.placeholder(dtype, shape=(None,1))`
- Step 2
 - Define tensorflow variable
 - `tf.get_variable(name, shape=(2,1), initializer=?)`
 - `tf.random_normal_initializer()`
- Step 3
 - Define linear model
 - `tf.multiply(x_ph,a_hat) + b_hat`
- Step 4
 - Define loss function
 - `0.5 * tf.reduce_mean(tf.square(y_pred - y_ph))`
 - Define optimizer
- Step 5
 - Run optimization

Tensorflow Graph Construction

Initialize graph

Train it

```
: tf.reset_default_graph()

x_ph =
y_ph =

a_hat =
b_hat =

y_pred =

loss =

# define optimizer
optimizer =
train =

# initialize variables
init = tf.global_variables_initializer()
```

```
nepoch = 10000
feed_dict = {x_ph:y_ph}
for epoch in range(nepoch):
    loss_np,_ =
    if (epoch%1000) == 0:
        print("[%d/%d] loss : %f"%(epoch,nepoch, loss_np))
```

[Exercise 3] General Linear Regression

- Linear model: $h_{\mathbf{w}}(\mathbf{x}_j) = \sum_{i=0}^n w_i x_{j,i} = \mathbf{w}^T \mathbf{x}_j$.
- General linear model:

$$h_{\mathbf{w}}(\mathbf{x}_j) = \sum_{i=0}^n w_i \phi_i(\mathbf{x}_j) = \mathbf{w}^T \phi(\mathbf{x}_j),$$


where $\phi(\mathbf{x}_j) = [1 \ \phi_1(\mathbf{x}_j) \cdots \phi_n(\mathbf{x}_j)]^T$ and $\phi_i : \mathbb{R}^n \rightarrow \mathbb{R}$ are nonlinear functions.

- Examples

$$\phi_i(x) = x^i \quad (\text{polynomial regression})$$

$$\phi_i(\mathbf{x}) = \exp\left(-\frac{1}{2s^2}(\mathbf{x} - \mu_i)^2\right),$$

where s and μ_i are fixed parameters.



Non-linear 1D data generation

```
def f3(x):  
    return np.sinc(x)  
  
x = np.linspace(-2*np.pi, 2*np.pi, 100)  
y = f3(x)  
  
n = 500  
x_data = np.linspace(-2*np.pi, 2*np.pi, n) + 0.05*np.random.randn(n)  
y_data = f3(x_data) + 0.05*np.random.randn(n)  
  
plt.plot(x, y, "k-", label="Ground Truth")  
plt.plot(x_data, y_data, "r.", label="Data")  
plt.legend()  
plt.xlim([-2*np.pi, 2*np.pi])  
plt.ylim([-1, 1.5])  
plt.show()
```

[Exercise 3] General Linear Regression

- General linear model: $h_{\mathbf{w}}(\mathbf{x}_j) = \mathbf{w}^T \phi(\mathbf{x}_j)$

- Let

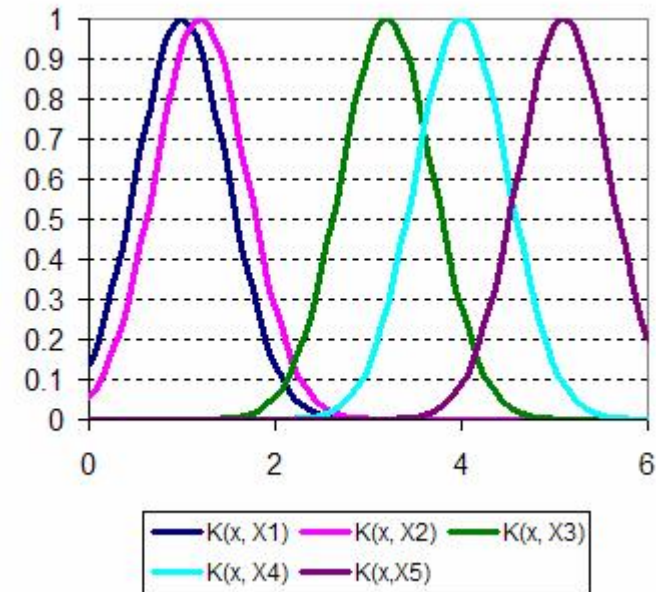
$$\Phi = \begin{bmatrix} 1 & \phi_1(\mathbf{x}_1) & \cdots & \phi_n(\mathbf{x}_1) \\ 1 & \phi_1(\mathbf{x}_2) & \cdots & \phi_n(\mathbf{x}_2) \\ \vdots & \vdots & \ddots & \vdots \\ 1 & \phi_1(\mathbf{x}_N) & \cdots & \phi_n(\mathbf{x}_N) \end{bmatrix}.$$

- Then

$$Loss(h_{\mathbf{w}}) = \sum_{j=1}^N (y_j - \mathbf{w}^T \phi(\mathbf{x}_j))^2 = (\mathbf{y} - \Phi \mathbf{w})^T (\mathbf{y} - \Phi \mathbf{w}).$$

- Now, we use matrix form!
- Goal
 - Construct the matrix Φ in tensorflow graph
 - Define predictor and loss function
 - Optimize weight parameters

Example of $\phi_i(x)$



$$\phi_i(\mathbf{x}) = \exp\left(-\frac{1}{2s^2}(\mathbf{x} - \mu_i)^2\right)$$

[Exercise 3] General Linear Regression

- The matrix Φ

Compute this first!

$$\Phi = \begin{bmatrix} 1 & \phi_1(\mathbf{x}_1) & \cdots & \phi_n(\mathbf{x}_1) \\ 1 & \phi_1(\mathbf{x}_2) & \cdots & \phi_n(\mathbf{x}_2) \\ \vdots & \vdots & \ddots & \vdots \\ 1 & \phi_1(\mathbf{x}_N) & \cdots & \phi_n(\mathbf{x}_N) \end{bmatrix}$$


The number of data : N

The number of kernel function $\phi_i : n$

- All elements look like $\phi_i(x_j) = \exp\left(-\frac{(x_j - \mu_i)^2}{2s^2}\right)$
 - Compute distance matrix
 - $D_{ij} = [(x_j - \mu_i)^2]$
 - Compute kernel matrix using element-wise operation
 - $\Phi = \left[1, \exp\left(-\frac{D}{2s^2}\right) \right]$

[Exercise 3] General Linear Regression

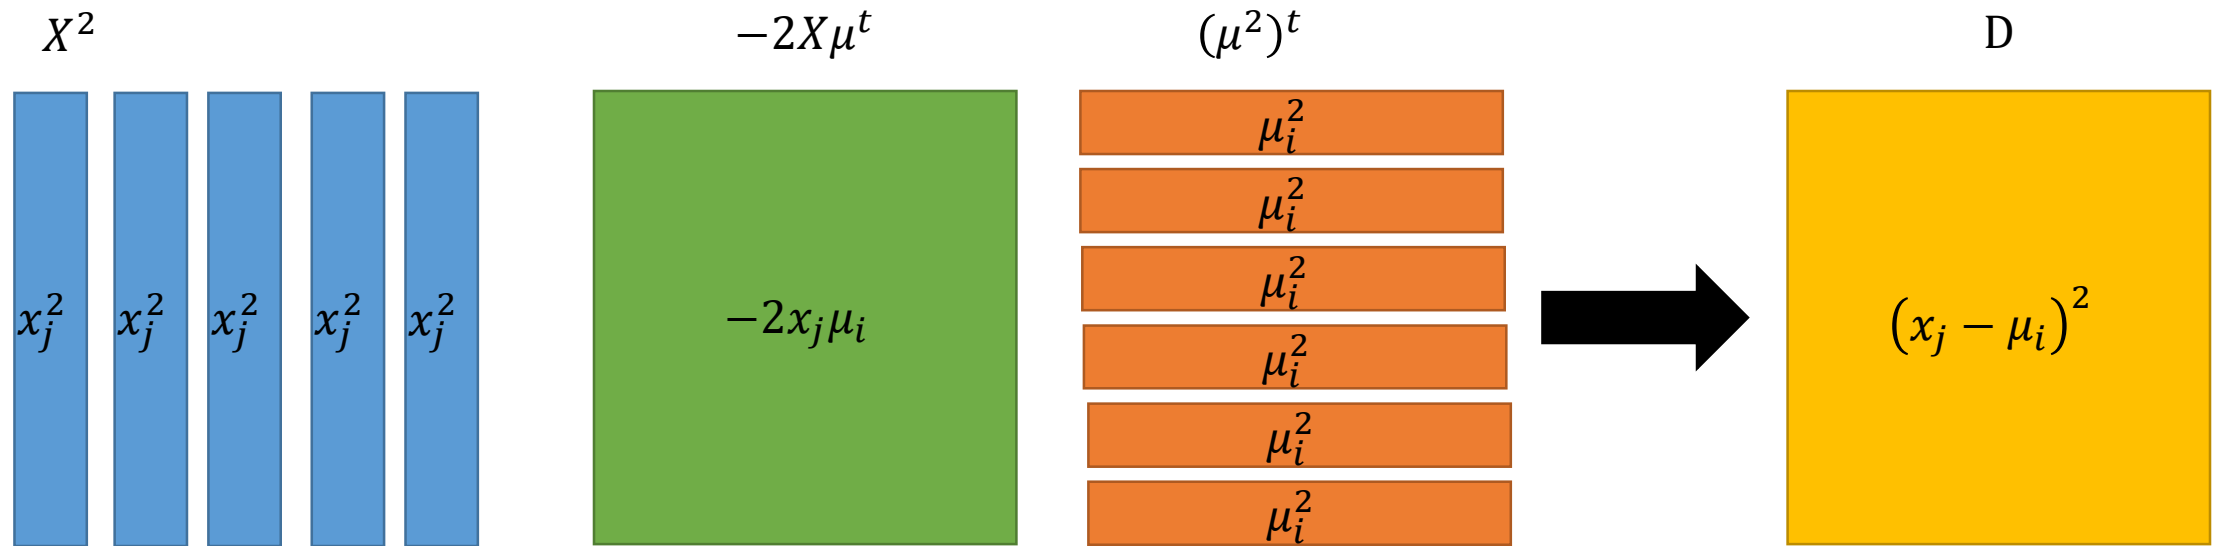
- The matrix D



- Python broadcasting!
- TF broadcasting!

[Exercise 3] General Linear Regression

- The matrix D



- Python broadcasting!
- TF broadcasting!

[Exercise 3] General Linear Regression

- Step 1
 - Define placeholder
 - x_{ph} : shape (None,1)
 - y_{ph} : shape (None,1)
 - μ_{ph} : μ_i
 - shape (n_kernel+1, 1)
 - $inv_squared_s_{ph}$: $h = s^{-2}$
 - shape ()
 - w_hat : w_i
 - shape (n_kernel+1, 1)
- Step 2
 - Compute X^2 : `tf.reduce_sum(tf.square(x_ph), axis=1)`
 - Compute μ^2
 - Reshape X^2 : `tf.reshape(x_norm, [-1, 1])`
 - Reshape μ^2
 - $D = X^2 - 2X\mu^t + \mu^2$: `tf.matmul(x_ph, mu_ph, False, True)`

Tensorflow Graph Construction

```
tf.reset_default_graph()

n_kernel = 20
mu = np.linspace(-2*np.pi, 2*np.pi, n_kernel)
inv_squared_s = 1e1

x_ph =
y_ph =

mu_ph =
inv_squared_s_ph =

w_hat =
print("Define place holders")
```

Distance Matrix D

$$X^2 - 2X\mu^t + \mu^2$$

```
x_norm =
mu_norm =

x_norm =
mu_norm =

squared_dist =
```


[Exercise 3] General Linear Regression

- Step 3
 - Compute kernel matrix
 - $\exp\left(-\frac{D}{2s^2}\right)$
 - Concatenate one vector
 - $\Phi = \left[1, \exp\left(-\frac{D}{2s^2}\right) \right]$
- Step 4

Train

```
init = tf.global_variables_initializer()

sess = tf.Session()
sess.run(init)

w_np = sess.run(w_hat)
print("Randomly initialized parameters")
print("w: {}".format(w_np.flatten()))

nepoch = 10000
feed_dict = {x_ph:x_data[:,np.newaxis],y_ph:y_data[:,np.newaxis],mu_ph:mu[:,np.newaxis],inv_squared_s_ph:inv_squared_s}

for epoch in range(nepoch):
    loss_np,_ = sess.run([loss,train],feed_dict=feed_dict)
    if (epoch%1000) == 0:
        print("%d/%d loss : %f"%(epoch,nepoch,loss_np))
```

```
kernel =
kernel =

y_pred = tf.matmul(kernel,w_hat)

loss = 0.5 * tf.reduce_mean( tf.square(y_pred - y_ph) )

optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.1)
train = optimizer.minimize(loss)
```

